Map of Denmark
First-Year Project, Bachelor in Software Development,
IT Univ. of Copenhagen

Group 12
Jakob Melnyk jmel@itu.dk
Niklas Hansen nikl@itu.dk
Emil Juul Jacobsen ejuu@itu.dk
Jens Dahl Møllerhøj jdmo@itu.dk

Supervisor: Lars Birkedal
Advisors: Jonas Brabrand Jensen and Filip Sieczkowski

May 25th, 2011

Contents

Chapter 1

Preface

Chapter 2

Background

2.1 Problem area

Over the last decade people have switched from traditional roadmaps to using the web-maps. This is a change without any negative side-effects. The online services remove all the problems with determining the quickest route between two points and you spend no time browsing the pages of the map to find what you need. With the popular smartphones the online map is even more useful, since you no longer need to prepare your trip before you leave.

The online maps have now been used for many years and haven't been slow at adopting new features to improve their usability. They have both implemented satellite-maps that allow us to browse the entire planet from above, and lately the feature called Google Street View has upped the stakes when allowing us to look at any direction from a given point of a road. The two maps that we use the most are Google Maps and the Danish map called Krak. These maps both have the mentioned features but slight differences in the way the user navigates and searches for routes.

Because of the widespread knowledge of the online maps, the users have been accustomed to certain features and ways of using the map. It is very important that we, with a new map program, use this knowledge to our advantage and don't try to reinvent the wheel. By using some of the commonly used controls in our map, a user will be able to quickly adapt to our program and use it efficiently.

2.2   Requirements for the map

Our teachers had a few requirements for features in the project. The requirements were presented to us during development, in 3 steps. This made it easier for us to focus on getting the basic features of the program to work, but it made it harder for us to plan ahead as well.

Here is a list of the features we had to implement:

- We had to make a visual representation of all the roads from the dataset.

- We had to draw different roads in different colors.

- We had to adjust our drawing of the roads to the windows size of our GUI.

- We had to make mouse zooming possible. The user should be able to drag the mouse from one corner of the map to another, in order to zoom in on the selected area.

- We had to implement a method to find a route from one road in the map to another, and we had to allow the user to find that route by clicking at each of the roads.

We had to consider these requirements while deciding how the program should be designed. Of course, some of the requirements were so self explanatory that we would have made them regardless of them being required.

2.3   Our requirements

In the process of designing the program we made some requirements that we wanted to make sure was met before making more advanced features. Since it was required to make the user able to zoom-in on the map, we found it logical to allow him or her to:

- Zoom out

- Scroll the map

With these basic features covered, we decided which advanced featured we wanted implement. In order to give the user a chance to find more specific places, we decided to show the user the name of the road closest to the cursor. This means that the user can get orientated without clicking or pressing any button.

The most interesting feature of our Map of Denmark is perhaps the option of selecting whether you want to travel by car or by bicycle. Many Danish people use bicycles to travel and as such this feature is a very relevant and nice one to have in the software. It is a feature not usually seen on international maps - this is most likely due to the size of these maps.

We also decided to let the user create routes with an unlimited amount of waypoints. This makes our map well suited for planning bicycling trips or longer car rides where you want to reach more than one destination. Because this feature can be quite demanding when a lot of waypoints are selected, we had to make sure that the algorithm for finding routes was fast enough so we avoided making use of the map cumbersome for an end-user.

We wanted to make sure that the shape of Denmark is recognizable when the user has zoomed out to show the entire country. There is a balance between showing a big amount of roads and the delay when navigating. We wanted to make the user able to navigate the map with a reasonable amount of delay.

The fact that we chose specific requirements for our program gave us two big advantages. It both made the planning process easier, and it made the structuring of the code easier to choose. In the process of creating the program, we had to change and create new requirements for ourselves, when we felt some feature was necessary for the end-user to have. In the last part of the coding process, we decided on our final requirements and worked towards completing them.

2.4 Data set

We have been provided with a dataset of roads and intersections in Denmark from Krak. Additionally we got some code for loading the data in from the text files. We have only made minor changes to the code for loading the data.

2.4.1 UTM-coordinates

It is important to note that the KrakNodes are in UTM-32 coordinates. When using the UTM standard the origo is placed at the south-west corner. These coordinates need some conversion when using in Java since the origo is placed differently.

2.4.2 Graph

When the data has been loaded it is stored as a Graph containing KrakNodes and KrakEdges. The KrakEdges are the road segments and contains the name of the road, an estimated drive time, a direction of traffic and references to

the two KrakNodes that are at either end of the road. The KrakNode itself contains only the coordinates for the point. The Graph itself contains a number of useful methods for searching the data like getting all edges that is connected to a KrakNode. We will be using these methods extensively throughout the project both for drawing the map and for finding the route between two points.

2.5 MVC structure

In order to achieve a decent code structure and separation of responsibilities, we utilize the Model-View-Controller (MVC) architecture. By doing this, we split the code into smaller sections, which can be easier to handle and easier to coorporate on the code.

This is because we split the code into three parts: graphical user interface ("view"), data handling ("model") 'and the connection between the two ("controller"). Read more in section 5.1 on page 21.

2.5.1 What is it? How does it work?

In order to achieve a decent code structure, it is important to split the responsibilities of the program into different pieces, which work together to make the program work as intended.

One way to do this, is to have the graphical user interface in a class of its own, and the rest of the program in another class. The downside to this approach, is that it can get ambiguously where to put specific pieces of code (button listeners and such).

We chose to utilize the Model-View-Controller (MVC) architecture, which is another way of structuring a program (a better one). With MVC we divide our code in three main parts, in order to achieve a decent separation of data, logic and the graphical interface.

Like the picture above shows, every graphical window has its own class. These classes are called "views". Where we beforehand had one class to handle both communication with the data and the graphical windows, we now split this into two: "models" and "controllers".

The models handle all communication with the data sources, and every model handle one data source. If we were to use relational databases, we would have a model for each table in the database. In our case we only have one data source (Krak's dataset), and therefore only have one model, to communicate with this.

Controllers Now we only need a way to connect the graphical interface with the data. This is where controllers enter the picture. In MVC you have one

Figure 2.1: Overview of MVC

controller per view. This controller provides the view with all the neccessary data from the models. The controllers also have listeners on the view, that listens to events on the view (like when the user presses a button, clicks his/her mouse and such). Then the controller can provide new data to the view or save new data from the view to the data source (through the models).

An example of this could be that the user updates some info and presses the "Save"-button. Then the controller listens to this event, passes the new data to the models and it gets saved. In our case we have listeners to (among other things) mouseclicks, so we can place pins when the user clicks the map.

Chapter 3

User Interface analysis

In this chapter we describe our decisions and present our analysis and arguments regarding some of the features that we find interesting.

3.1   User interface as a whole

When we designed the first version of the graphical user interface in the first part of the project, we decided to make a window inside of the graphical user interface where the actual map should be displayed. We chose to have this window placed on the right side of our graphical user interface and interaction with the user mainly placed on the left.

We believe that this is a simple way of representing a user interface for a map. A lot of software use a menu bar with dropdown menus for selecting different functions. When we designed our outline for the graphical user interface, we did not design it with a huge amount of functions in mind.

The features that we have implemented in this version can easily fit in our simple user interface, but if features like searching for roads, route planning, etc. are included, then space and overview may become an issue on the left side.

Depending on the feature, we feel it would be beneficial to let the main window change when different feature types are selected.

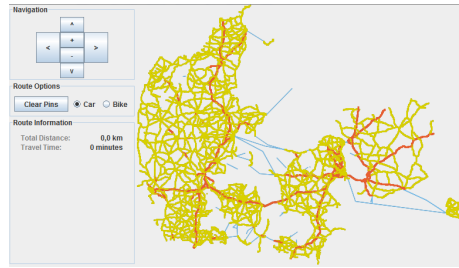Below is a screenshot of our standard user interface. How to use it will be explained in the Manual on page 24.

Figure 3.1: Screenshot of GUI

3.2 Interesting features

This section presents some of the interesting features we have implemented.

3.2.1 Zoom

We have a few options for zooming in and out on the map. As described in section 2.2 Requirements for the map on page 6, it was required that we made it possible to zoom by dragging a box around the part of the map the user wants to look closer at.

In addition to the option of using the mouse to zoom, we have implemented a zoom-in and out function on the GUI and a hotkey for zooming out to the original view. We made the original view function a hotkey only because we did not want to have too many buttons on the left side. We considered making it a menu bar function, but we did not manage to get it into this version.

We felt we really needed a zoom out function, so users do not need to close the program and start it again, when the user wants to view the map further zoomed out. A combination of the zoom in and out functions helps the user a lot when navigating the map.

We have limited how far a user can zoom in and out. If the user tries to zoom out further than the original zoom level, the view will default to the original zoom level. If the user attempts to zoom in further than a width or height of 200 in UTM32-coordinates, the zoom function will do nothing. This could probably be improved by zooming in on the smallest possible area at the position the user selected, instead of doing nothing.

3.2.2 Navigation

We have made it possible for the user to navigate the map by using the arrow buttons on the graphical user interface. When one of the buttons are pressed,

the "view" will move in the direction specified by the button. While it was not specified as a requirement for the project, we felt it was a necessity to implement at least basic navigation functionality.

Like we did with the two zoom functionalities, we have limited how far a user can move around the map. The user is free to move around the map, but if user moves outside the bounds of the map in a way where the view would show an image that is not part of the map, the move function will not do anything.

3.2.3   Hotkeys

We have implemented hotkeys for all the buttons on the graphical user interface plus an additional for zooming back to the original view. When we discussed the benefits of hotkeys, we felt it was important for experienced users of the software should have a less cumbersome time navigating the map.

At first we just had hotkeys for the clearing of markers (mentioned in section 3.2.4) and zooming out to the original view, but we later added the hotkeys for the rest of the functionalities. If more features are added in a future version, it would be important for us that a hotkey were provided, if at all possible.

3.2.4   Route planning and markers

Part of the requirements for the project was to provide the user with a way to get the fastest or shortest route from one point to another. We accomplish this by putting a "marker" at the spot where the user clicks with the mouse. The marker shows which number in the sequence of markers it is. This will change if a marker is removed. Originally we had "pins" instead of markers, but we changed it, as we felt the pins we had were a bit large.

We have made it possible to place more than the two markers the project requirements asks for. If the user places more than two markers, the software will find the shortest route between 1 -¿ 2 and 2 -¿ 3 and so on. This was cheap for us to implement, and we felt it added a nice touch to our program.

We have implemented two methods of removing pins from the map. We have assigned a hotkey to the graphical user interface button "Clear Markers", which removes all the markers from the map. The other way of removing pins is by clicking on them. This functionality is both intuitive and confusing at the same time. It is intuitive to click the marker you have just placed if you want to remove it, but it is not obvious in our interface. We believe that it is enough to have the "clear all markers" functionality for those who do not find it intuitive to click markers to remove them, and for the users that do find it intuitive, we offer them an easy way to undo a missclick.

3.2.5   Bike/car

Another interesting feature in our Map of Denmark project is the option to switch between bike and car routes. The user interface will start with car selected when the program starts.

If a route is marked by the user, it will display the length and the estimated travel time on the left part of the user interface. When the bike option is selected, it recalculates the route for a bike, without visiting highways and other roads that bikes cannot or are not allowed to drive on. If the user switches back to the car mode, it recalculates the route again, but not visiting small paths and other roads where a car is not allowed to drive. The estimated travel time is also recalculated. The user does not need to have planned a route before he/she changes the type of transportation.

We have implemented this to help our software target a wider group of people. The bike/car options were a bit costly to implement, but we categorized it as a very beneficial feature and we did not feel we could leave it out.

3.3   Features not implemented

This section presents some of the features we have chosen not to implement. These features are not in the final program, because we did not feel there were compelling arguments for implementing them.

Features that we wanted to implement, but did not make it into the final version, will be discussed in chapter Product conclusion on page 28.

3.3.1   Choice of roads to be displayed

We chose not to implement the option of selecting which roads to be displayed. In a sense our program already does this by showing more detail the further zoomed in the map is. It could become very confusing if the user had the option of selecting roads, because the graphical user interface could become very cluttered, if all the roads were listed.

One advantage to this could be the option for the user to select which roads should be included in the route planning - ferries, highways, bridges etc.

3.3.2   Smooth scrolling

We made an attempt to let the keyboard arrows scroll smoothly over the map, but we could not implement it to work fast enough, so the user would experience

lockups and the user interface hanging at some points. A solution to this would be to save map at images that you can scroll across - this would be faster, but would require more disk space. Because we store the data the way we do, which forces us to draw every line individually every time the user moves the viewport, we cannot do this fast enough.

In the end, we decided the benefit of the smooth scrolling was not big enough for us to spend a lot of time implementing this feature. The cost of changing that much way we draw the roads, was simply too high compared to the benefits.

3.3.3 Dynamic route finding

In the final project description, a dynamic route finding feature was suggested. We considered implementing this as we thought it was a nice feature to have, but it conflicted with the algorithm we use for calculating the route. More about the algorithms can be read in section 4.3 Dijkstra vs A-star on page 17.

Chapter 4

Implementation

This chapter describes how we have implemented some of the more interesting features of the software. We aim to describe it to enough detail that this chapter can serve as a guideline for implementing the functionalities we describe.

4.1   UTM-conversion

When the graphical user interface part of the map tries to communicate with the model through control, some conversions of the different kind of values are necessary. Both when going from coordinates in the java-coordinate system to UTM32-coordinates and back.

We need to convert the values when we want to use the mousezoom and when we want to place the markers for pathfinding. We get an input on the graphical user interface when we mousezoom and this needs to be converted to UTM-coordinates so that we can create the new boundaries of the zoomed rectangle.

When we place markers for pathfinding, we do the same as when we do mousezoom, but instead we store the point as UTM32-coordinates and whenever we move the map, we convert it back to pixel-coordinates so that we know where to draw.

The java-coordinates have origo in the top left corner with the y-coordinate increasing the further down the y-axis you go. UTM32-coordinates are a bit different. UTM32 has origo in the bottom left corner and the y-coordinate increasing the further up you go on the y-axis.

We have a utility class with methods for converting the points back and forth. One takes a point from the view, the model and the view itself uses this formula for converting the pixelpoint to the UTM32-point.

$$a_y = canvas_{height} - a_y$$

$$UTM_x = bounds_x + \frac{a_x}{canvas_{width}} \times bounds_{width}$$

$$UTM_y = bounds_y + \frac{a_y}{canvas_{height}} \times bounds_{height}$$

return $\mathsf{point}(\text{UTM}_x, \text{UTM}_y)$

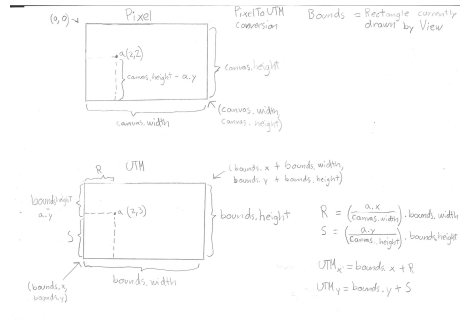Below is an illustration of the conversion from pixel to UTM.

Figure 4.1: UTMillustration

To convert from UTM to pixels we use the same formula but reversed.

$$Pixel_x = \frac{(a_x - bounds_x)}{bounds_{width}} \times canvas_{width}$$

$$Pixel_y = \frac{a_y - bounds_y}{bounds_{height}} \times canvas_{height}$$

$$Pixel_y = canvas_{height} - Pixel_y$$

return: $\mathsf{point}(\text{PIXEL}_x, \text{PIXEL}_y)$

4.2   Mousezoom

We have implemented the mouse zoom requirement by using $\mathsf{mouseEvents}$ on our canvas. When the user presses the left mouse button down, it generates a $\mathsf{mousePressed}$ event. We record the position of the mouse at the time of the $\mathsf{mousePressed}$ event and wait for the $\mathsf{mouseReleased}$ event. When the left mouse button is released, it generates a $\mathsf{mouseReleased}$ event. We use the position of

the mouse at the **mouseReleased** event and the **mousePressed** to calculate new bounds for the model.

As mentioned in section 4.1 **UTM-conversion**, we use UTM-conversion to convert pixel values into the UTM values we need.

Often the user will not drag a square that is in perfect ratio with the canvas. If it does not have the same ratio as the canvas, we change the ratio of the dragged square behind the scenes. We do this by either adding length or width to the dragged square. We always make sure to at least show what was inside the box the user dragged.

If the ratio of the dragged square is smaller than the canvas ratio, we make the dragged box wider. If the ratio of the dragged square is larger, we make the dragged box taller.

4.3 Dijkstra vs A-star

In preparations to implementing the path finding feature to our program, we knew two possible choices. They are named Dijkstra and A* (A-star), and are quite similar but behave a little different.

The Dijkstra algorithm uses a minimum priority queue to find the shortest path from a given node to every other node by looking at the edges connecting nodes. The program will take a node from the priority queue and add all the other nodes that are connected from the current to the priority queue. The priority queue takes a value to the node and this should be the distance to the current node plus the length of the edge between the two. Since the priority is made to return the node with the smallest value associated with it, the next node in line will always be the one which is closest to the start node. This procedure continues until all nodes have been visited and by logging what edge led to all the nodes it is possible to trace back the route to the start node.

This algorithm is great if need to find the distance from one point to many other points, but can be quite slow since it just searches in all directions without concerns to the direction of the target node.

This is where A* comes in handy. The A* algorithm is a modification to Dijkstra that also looks at the estimated distance from the given node when determining the value for the priority queue. When using the geographical distance as a measure of best route the value would be the current distance from the start node plus the direct distance to the target (as if there were a road directly to the target). With this subtle change the algorithm will prioritize nodes that are relatively closer to the target rather than those that are in the other direction. This makes the algorithm much faster since it will not pay much attention to the roads that are not in the direction of the target. We have decided to use

the A* algorithm since we only calculate routes between two distinct nodes and therefore don't need the route from the start node to all others. The time reduction that A* gives is also a definite plus since no user wants to sit and wait too long for the program to find the route.

4.4   Evaluator

In order to make our path finding algorithm flexible enough for different interpretations of the "best route", we have added an entity called Evaluator. This is an object that has the responsibility of evaluating a node relative to the target node. The Evaluator also has the responsibility of calculating the heuristics that the A* algorithm relies on. By using the Evaluator we are able to use the same path finding algorithm for two very different tasks, namely the biking route and the car route. The major difference between these is that the bike uses the distance and the car uses the total drive time. This implementation is also a good example of making our code ready for future features, since if we needed to add other means of transportation or simply variations of the ones we have, we would only need to create new Evaluator objects and not change a single line of code in the A* algorithm.

4.5   Quadtree

In order to improve the drawing of our map, we have implemented the data structure Quadtree. The Quadtree divide our road data into smaller rectangles. When we want to retrieve data from the map, we can give the quadtree a rectangle, and it will return all the roads within rectangles that intersect that rectangle. This technique optimizes the drawing of the roads, because we donÕt draw roads outside of the view. However, when viewing the entire map of Denmark, this implementation does not help us. Therefore, it is necessary to only to draw the bigger roads when zoomed out. We have discussed two different techniques to do this. The first technique relayed on putting the bigger roads at the top of the quadtree when building it. Then we could specify at which depth we wanted to search the quadtree. The second technique relayed on dividing the quadtree into different parts, each containing their own road types. Going with this implementation, we simply specify which roads we want to draw. This method requires more RAM than the first one, because it requires more instances of object QuadTreeNode. At first glance, the first method seemed to be superior to the second, because it saves some RAM. However, the second implementation is easier to implement, and the extra RAM are such a small percentage of the overall RAM use, that it does not really matter.

4.6 Serialization

We observed that the user had to wait quite a long time for the program to start. This was because every time we start the program, we loop through the entire dataset given to us by Krak. This data-set is huge, and because of this it takes quite a lot of time to start the program. Because we need to load all these data, the user is presented with a blank screen for a long time, before all these data are loaded and the program starts.

We started looking for a way to speed up the loading process, so the user has a map in front of him or her quickly, when the user starts the program. What takes the most time is looping through the data and creating the needed datastructures (quadtrees, the graph and so on), so if we could skip these steps or speed them up, we could save a lot of time.

This is where serialization comes into play. By serializing an object, you transform your object into something that can be passed around, through streams and such. So by serializing objects, you can save them to files. If the object to be serialized contains references to other objects, these will also be serialized (if they are Serializable / implements java.io.Serializable).

By doing this, we only need to build our datastructures the first time you start the program. After the objects have been created, they are been serialized and saved to files. The next time the user starts the program, we check whether the data has been changed. We check this by checking the MD5 checksum of the file, with MD5Checksum (we didn't write this class ourselves). When we serialize the objects, we also save the checksum to a config file. If the data hasn't been changed (i.e. the checksum is the same), we load the objects that we saved, instead of making them all over.

We serialize all the quadtrees, the graph and the maxbounds-object (specifying the bounds of Denmark), as these are the objects we need for the program to run. When saved to the harddrive, it is around 65MiB, which is okay, given the sizes of harddrives today.

We save all these objects to one single file, in order to keep the references. If we didn't do this, the references will be ruined, which can break the program. We experienced problems with this, as nodes and edges is both stored in the quadtrees and the graph. If we serialized and saved the quadtrees in one file, and the graph in other file, a given node will be saved twice, and when we load the data from the serialized files, the node will exist twice, and it won't be the same node. But if we save the objects to a file through the same stream, we only get one of each, which leads to less RAM usage, less harddrive usage, faster loadtimes and fewer bugs.

Threading

By serializing our main objects, we cut several seconds of our load times. But we can do it even faster. We are serializing several objects, but only few of them are needed right from the start of the program. So what we can do to speed it up even more, is loading the few necessary objects, and then load the rest in the background. We do this with threads. We load the few objects we need from the start, then create a new thread to load the rest of the objects, and in the mean time, we create the window and draw the map.

The same goes for the first run. The user doesn't need to wait for the program to finish serializing and saving to files. By using threads, we can create the datastructures, and then immediately show the window to the user, while saving the objects to files in the background.

Chapter 5

UML-diagrams

5.1 MVC

Figure 5.1: UML of the MVC architecture in our Map program

Figure 5.1 shows our implementation of the MVC architecture. Because we have exactly one window, we chose to name our view and controller "View" and "Control", respectively. The same goes for the models. We only have one data source (Krak's data-set), og because of that our model is called "Model". We could have named these three classes something different that may have been more meaningful in terms of the Map of Denmark context, but we chose these name in order to make our architecture clear.

The easiest way to understand MVC and how the individual parts talk together, is by using an example. Let us say the user has already clicked the map and

placed a pin. Now the user clicks on the map again to place another pin.

The controller has placed listeners in the view, so when a MouseClicked event is thrown, the controller is called. First it checks if there is another pin at the spot of the click. If there is, this will be removed, and the model is told to clear the route. If there is still over two pins placed, the model is asked to calculate a new route.

If there isn't a pin where the user clicks, we place a new pin. If the user has placed two or more pins, the controller calls its own findPath-method from point 1 to point 2, point 2 to point 3 and so on. The findPath()-method tells the model to find a path between the two points given.

The model then asks a helper-class to find a path, using the A* algorithm, and provides it with the graph and the two points. When a path is found, it is saved in the model, ready for use in the controller.

The final step is getting the view to draw the route. The controller gets ready for a repaint, by fetching the route from the model. Then it passes this route to the view's repaint-method, and the view paints the road.

5.2 Simple Diagram

5.3 Control flow

Chapter 6

Tests

6.1 WhiteBox: closestEdge

6.2 JUnit

The model classes have some of the most interesting functions and algorithms of our program. The most advanced functions in the model are those regarding path finding. The task of finding a path from one node of the graph to another, takes many different kinds of input. Here is a list of the tests we ran on the model class: The tests are chosen so that they test for different input. But they are also chosen so that they test the logical problems our path finding algoritm might run into. In order to test these problems proberly, we constructed some simple testdata that we knew we could relay on: Here is shown a visualization of the graph our testdata constructed. The graph is constructed so that we can test the many different problems that might occur on a small and reliable data set.

6.3 System test

Chapter 7

Manual

7.1 Navigation

At startup, the entire country of Denmark is shown. Now, we can move around the map with both the diretion buttons in the top left corner, and the arrow keys. To move west, click the button $\boxed{<}$ or press the left arrow key. This goes for all 4 directions.


Figure 7.1: Zooming in on Copenhagen

7.2 Zoom

To zoom-in on the map, click the ⬭ + ⬭ button in the navigation panel.

To zoom-in on a specific area, click and drag a rectangle around that area on the map. A blue transparent rectangle will show you the selected area. To zoom-in, release the mousebutton.

Fore example, if you want to zoom-in on Copenhagen, click the upper left corner of the city, and drag the cursor to the lower right corner.
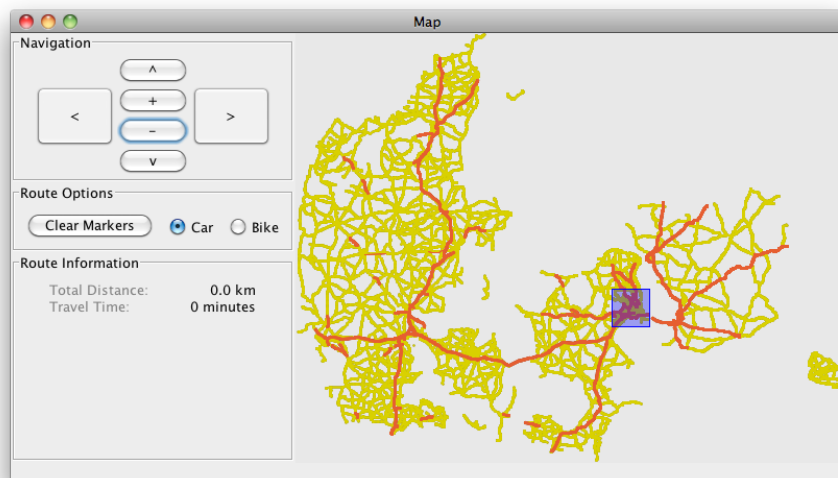

Figure 7.2: Zooming in on Copenhagen

To zoom-out of the map, click the ⬭ - ⬭ button. To return to the startup view (showing the entire map of Denmark) press the esc button.

7.3 Route find

To find a route from one point to another, you must specify a start and an end location. Click anywhere on the map to choose your start location. A light blue marker will appear, containing the number Ò1Ó. To choose your end location, click at another location. A marker containing the number Ò2Ó will appear. The best route from 1 to 2 will be calculated, and shown on the map as a blue path. To extend your route with more markers, you can click at a new location on the map. You can repeat this an unlimited amount of times. You can delete

one of your markers by clicking at its root. To delete all marker, click the button 'clear markers'.
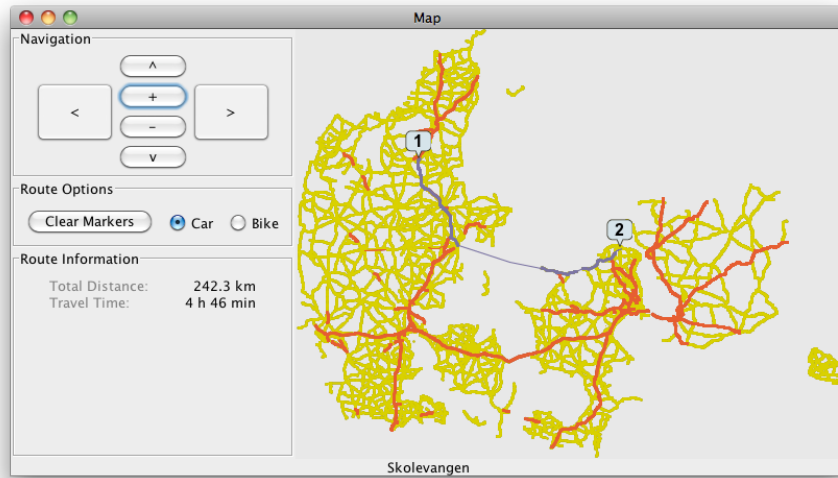
Figure 7.3: A route from 1 to 2 has been calculated

7.4 Bike/car

When calculating routes, it is important to specific which form of transportation you wish to use. Choose your preferred style of transportation by selecting the corresponding radio button.

Figure 7.4: The radiobuttons

7.5 Resize

To resize the map, drag the window as you would with any other application. The map will automatic adjust to the new size of the window.

Figure 7.5: The map has been resized

Chapter 8

Product conclusion

Chapter 9

Group norms

We wrote a constitution before we the project was handed out to us. In this constitution we describe what we require of eachother and ourselves in terms of working on the project. We felt we made some rather strict requirements so that we were sure to get some work done. This backfired a little bit, because we did not keep all the agreements we made, but most of the worked and we later changed them a bit once our schedule cleared up from lectures early May.

This is the requirement part of our constitution:

- Check mail at least once a day
- Tell the rest of the group in time if you have trouble getting done on time.
- Admit when you are not done on time.
- Do not waste time when we have meets.
- Respect that different people work in different paces and different ways.
- We need to evaluate often.

We also tried to get a mean of our level of ambition. Our goal was always to do what we could manage to do in the time frame that we had and without wasting time.

9.1   Meetings

We structured our work together in "meetings". A meeting was whenever we were together working on the project - these were to be done at the ITU. The structure of a meeting was simple:

- Leader of the meeting presents his plans, if any, he has for today.

- Leader selects someone to write down what happened at the meeting.

- What have we done since last time?

- Who does what today?

- Work today.

- Fifteen minutes before work ends: decide on homework for next time and select leader of the next meeting.

Before our lectures ended in early May, we had meets Tuesdays and Fridays from 10 AM to 4 PM. After lectures ended, we felt only meeting Tuesdays and Fridays would be too little time spent in meetings. So we decided to make our meetings one and a half hours shorter, but instead meet on Mondays, Wednesdays and Fridays from 10 AM to 2:30 PM.

Chapter 10

Diary

In this project we have kept two types of diaries. One type in our meeting documents and one in a spreadsheet. The meeting document diary was done from meeting to meeting, so that we kept track of what everyone had done before each meeting.

The spreadsheet diary was kept as a seperate document where we wrote down whenever we had done something outside of meetings. We were not so dedicated to writing down everything we did in this diary, so it ended up being incomplete. We have included it anyway though, as we feel it has been a part of the process.

Our worksheets can be found in appendix B on page 39. They are written in Danish, as we had all of our meetings in Danish. The spreadsheet diary is in appendix B.1 Spreadsheet diary on page 39. The spreadsheet diary is also written in Danish.

Chapter 11

Worksheets

Chapter 12

Process description and reflection

Bibliography

[1] Robert Sedgewick and Kevin Wayne. *Algorithms, Fourth Edition*. Preliminary Edition Fall 2010. Addision-Wesley 2010.

Appendix A

Tests

This section covers tables and other appendices for our Tests chapter in the report.

A.1 PointMethods test table

The table in this section shows the full extent of our JUnit test coverage of the PointMethods class. Any variable that is described in the example column in the row of the method name is used in all the tests for that method.

Test description	Example	Expected Output	Actual Output
testPixelToUTM	view:(0, 0, 800, 600) model:(0, 0, 4000, 3000) point(80, 350)	(400, 1250)	(400, 1250)
testUTMToPixel	view:(0, 0, 800, 600) model:(0, 0, 4000, 3000) point:(pixelToUTM(80, 350))	(80, 350)	(80, 350)
testPointOutOfBounds	view:(0, 0, 800, 600)		
Inside	Point a:(100, 400)	a = (100, 400)	(100, 400)
Far left	Point a:(-100, 400)	a = (0, 400)	a = (0, 400)
Far right	Point a:(1200, 400)	a = (800, 400)	a = (800, 400)
Far up	Point a:(100, -300)	a = (100, 0)	a = (100, 0)
Far down	Point a:(100, 2000)	a = (100, 600)	a = (100, 600)
Far left & far south	Point a:(-100, -200)	a = (0, 0)	a = (0, 0)

A.2   RectangleMethods test table

The table in this section shows the full extent of our JUnit test coverage of the
RectangleMethods class. Any variable that is described in the example column
in the row of the method name is used in all the tests for that method.

Test description	Example	Expected Output	Actual Output
testNewBounds	old:(50, 50, 100, 100))		
Direction.WEST	length: 0.5 Direction.WEST	(0, 50, 100, 100)	(0, 50, 100, 100)
Direction.EAST	length: 0.5 Direction.EAST	(100, 50, 100, 100)	(100, 50, 100, 100)
Direction.NORTH	length: 0.5 Direction.NORTH	(50, 100, 100, 100)	(50, 100, 100, 100)
Direction.WEST	length: 0.5 Direction.WEST	(50, 0, 100, 100)	(0, 0, 100, 100)
Direction.OUT	length: 0.1 Direction.OUT	(40, 40, 120, 120)	(0, 50, 100, 100)
Direction.IN	length: 0.1 Direction.IN	(60, 60, 80, 80)	(60, 60, 80, 80)
testFixByInner-Rectangle	a:(50, 50, 173.94, 100) b:(50, 50, 139.45, 100)	Absulate value of (Ratio of a - ratio of b) less than 1e-9	Value is less than 1e-9.
testFixByOuter-Rectangle	a:(50, 50, 633.51, 100) b:(50, 50, 293.57, 100)	Absulate value of (Ratio of a - ratio of b) less than 1e-9	Value is less than 1e-9.
Mousezoom	model:(0, 0, 1000, 750) view:(0, 0, 800, 600)		
Standard mousezoom	a:(80, 540) b:(480, 240)	(100, 75, 600, 375)	(100, 75, 600, 375)
Zoom too far on x-axis	a:(130, 5) b:(140, 150)	(0, 0, 1000, 750)	(0, 0, 1000, 750)
Zoom too far on y-axis	a:(130, 442) b:(340, 432)	(0, 0, 1000, 750)	(0, 0, 1000, 750)
testPoint2DToRectangle			
a.x<b.x && b.y<a.y	a:(50, 150) b:(150, 50)	(50, 50, 100, 100)	(50, 50, 100, 100)
a.x<b.x && b.y<a.y	a:(50, 150) b:(150, 50)	(50, 50, 100, 100)	(50, 50, 100, 100)
a.x<b.x && b.y<a.y	a:(50, 150) b:(150, 50)	(50, 50, 100, 100)	(50, 50, 100, 100)

a.x<b.x && b.y<a.y	a:(50, 150) b:(150, 50)	(50, 50, 100, 100)	(50, 50, 100, 100)
a.x<b.x && b.y<a.y	a:(50, 150) b:(150, 50)	(50, 50, 100, 100)	(50, 50, 100, 100)

A.3 System tests

Tests

- Navigation Buttons

    - Zoom in/out
    - Pan(north, south, east, west)

- Boundaries

    - Minimum zoom
    - Maximum zoom
    - Pan out of map(north, south, east, west)

- Hotkeys

    - ESCAPE (reset to full view)
    - C (clear all markers)
    - I (zoom in)
    - O (zoom out)
    - Arrows (up, down, left, right)

- Pathfinding

    - Place marker
    - Two markers
    - Two markers (same node)
    - Multiple markers
    - Clear markers (screen button)
    - Switch between modes (recalculate route)
    - Closest node (that fits current mode)
    - Cannot travel opposite on one-way streets

- Resizing

- Narrower (add to the top and bottom)
- Wider (cut top and bottom)
- Shorter (add to the sides)
- Taller (cut sides)

- Mouse zoom

  - Same ratio(show exactly what is marked)
  - High rectangle(add extra to fit ratio)
  - Wide rectangle(add extra to fit ratio)

Appendix B

Worksheets

B.1 Spreadsheet diary

This appendix covers the diary we made in spreadsheet format in the time working on this project. It is not a complete diary, as we did not always keep it up to date.

Dato	Medlem	Løst opgave	Personlig note
3/8/ 2011	Jakob	Oprettede dokument	
3/14/ 2011	Niklas	Lavede .gitignore, refaktorede kode(+prettify)	
3/15/ 2011	Emil	Implementeret QuadTree	Der er noget galt i Map, da op/ned scroll virker omvendt, desuden en mystisk implementation i Direction Enum'et, det skal ændres til en switch direkte i Map.
3/17/ 2011	Jakob	Fixede merge problems i control	Forstår ikke helt det med file. Skal have det forklaret på mødet. Skal bede Niklas om ikke at lave min hjemmeopgave for mig! :D
3/22/ 2011	Jakob	Lavede noget på mousezoom	
3/24/ 2011	Emil	Vejnavn ved klik	Er stadig nogle bugs der skal fixes

3/25/ 2011	Emil	Resize og thickness	
3/27/ 2011	Emil	Vejnavn...	Der står nu ingen tekst nederst når man har musen over 200m fra en vej, desuden ser firkanten en del federe ud (gennemsigtig blå) når man er ved at zoome
3/28/ 2011	Jakob	Testede på hjemmecomputer	Begyndte at kigge på LateX. Mulighed for at zoome ud til start ville være bonus mega awesome!
3/28/ 2011	Niklas	Optimerede RAM forbrug	Double ->Float
3/31/ 2011	Jakob	Tegnede hvordan vi laver pixels om til UTM og skrev det meste af det jeg skulle skrive	Vi har nem mulighed for at lave god formatering af vores tekst - skal huske at vise hvordan fredag. Ville være fantastisk, hvis vi kunne få lagt tingene ind i ordenlige packages - so much shit in default package! Diskuter PixelToUTM e.x/width*width
3/31/ 2011	Filip	Skrev afsnit om Line og tilføjede tykkelser på veje i map	
4/5/ 2011	Emil	refactoring	Flyttet noget logik fra QuadTree til Model
4/7/ 2011	Jakob	Refactorede control og oprettede nye klasser - lavede JUnit	
4/12/ 2011	Niklas	Tilføjede multithreading loading af serialiserede filer.	Langsomme computere kan lave forespørgelserne i nogle af quadtrees før de er loadet. UPDATE d. 15/4/2011 - det er fixet, ved at returnere de edges og nodes som er fundet i de fundne quadtrees.
4/23/ 2011	Jakob	Lavede clearpins	

4/28/ 2011	Niklas	Lavede en default evaluator	Hidtil var der ikke nogen bestemmelse om hvad der skulle ske, hvis der ikke var valgt en evaluator. Jeg har nu lavet en Evaluator.DEFAULT, der bare refererer til Evaluator.CAR. Hvad skal der ske hvis brugeren skifter evaluator, efter ruten er lagt?
4/30/ 2011	Niklas	Fixede serialisering	Det hjalp at serialisere til samme fil, med samme stream. På den måde opstår der ikke dublikater. Desuden er fylder dataen mindre, det går hurtigere, og bruger væsentligt mindre RAM (550MiB =>260MiB)
5/1/ 2011	Niklas	Forårs-rengøring af kode	Har flyttet klasser rundt til de rigtige pakker. Har desuden gennemgået koden sammen med FindBugs, og ryddet op i koden.
5/5/ 2011	Niklas	Tilføjede anti-aliasing	Problemet var at vi tilføjede anti-aliasing på baggrunden også, og derfor kørte det langsomt. Prøvede kun at bruge det på stregerne, og det fungerer ganske udemærket, og er stadig hurtigt.
5/6/ 2011	Emil	Skønhed	Implementerede "tykke" veje med outline og et mere differencieret quadtree PS: jeg har glemt at skrive herinde en masse gange
5/8/ 2011	Jakob	LaTeX	Arbejdede med LaTeX og med smooth scrolling
5/14/ 2011	Emil	White Box Test	Har skrevet white box tests til getClosestNode metoden, mangler dog at lave det til JUnit da jeg skal vide hvordan jeg loader test-grafen, de andre tests der bruger den virker ikke ved mig